

Calibrated Language – words that mean the same thing every time

4 July 2026 · for Heinrich's sign-off · extends the locked Redesign Spec (2026-07-03), changes none of it

The problem being fixed (your dad's feedback): Verbatim's findings are written in words, and the words are uncalibrated twice over. The AI picks "very positive" vs "positive" by mood, not by rule – the same data could be worded differently on two runs. And even a consistent word gets read differently by different people. (This is a famous failure: a 1960s CIA study found readers interpreted "probable" as anything from a 25% to a 90% chance.)

Proof it's real, from our own data – in the 3 July Ossur run, the model's "judgment" scores came out 8–9 confidence and 7–10 opportunity on every insight, and **all four recommendations were rated "high" priority**. When everything is high, the word high means nothing.

The fix, in one line: code rates, AI explains. Every "how much / how strong / how sure" word on a page is assigned by a fixed rule from measured data – never chosen by the model. The model's prose keeps doing what it's good at: what people are saying, why it matters, what to do. Magnitude lives in badges; meaning lives in sentences; they never mix.

This splits every claim into two kinds, each with its own treatment:

Kind of claim	Examples	Treatment
Measured fact – we counted it, because we read every comment	How prevalent a theme is; sentiment split; number of voices	Shown as a fixed word <i>with the count next to it</i> : Dominant · 21 of 36 Ossur conversations
Model judgment – the AI's opinion	Confidence, opportunity, priority	Never shown as a number (spec rule, unchanged). Expressed only through the existing evidence tiers, forced ranking, and sort order.

The counts are also a sales weapon: competitors sample, so they can't print real denominators. We can put "of 2,651 comments read" under every badge.

Ladder 1 · Prevalence – *how much of the conversation is this?*

Applies to themes. Computed from two things we already store per theme: distinct supporting **conversations** (videos) and distinct **voices** (comments). The share is measured against the theme's own group – Ossur themes against Ossur's analysed conversations, Ottobock's against Ottobock's – because group sizes differ 5× and a raw percentage would bury client themes.

Badge	Rule (config, tunable)	Real example from the 3 Jul run
Dominant	≥ 40% of the group's conversations, min 10	"Questions about Ossur products" – 21 of 36 Ossur conversations (58%), 54 voices

Widespread	≥ 15% of the group, min 5	"Inspired by athletic perseverance" – 8 of 36 (22%), 26 voices
Recurring	2+ conversations, below Widespread	"Questions about phantom pain" – 3 of 184 industry conversations
Early signal	1 conversation (existing tier, unchanged)	"Exhausting insurance approval battles" – 1 conversation

Badge always carries its evidence: **Dominant · 21 of 36 Ossur conversations · 54 voices**. Same word, same meaning, every page, every week.

Ladder 2 • Sentiment – *how do they feel about it?*

Computed from the measured sentiment split over rated conversations (already in `run_summary`). Note the shape-words: **Polarized** exists so strong-both-ways data is never mislabeled "mixed" – this is the escape hatch for your "vocabulary can't represent the situation" worry.

Badge	Rule (pos/neg = share of rated conversations)
Strongly positive	positive ≥ 70% and negative ≤ 10% (<i>the current run: 89% positive of 255 rated → this</i>)
Leaning positive	positive leads negative by ≥ 15 points
Balanced	neither side leads by 15 points
Polarized	positive ≥ 30% <i>and</i> negative ≥ 30% (checked first)
Leaning negative / Strongly negative	mirrors of the positive rungs

Ladder 3 • Evidence strength – *unchanged, just named*

The spec's curation gate already is a calibrated ladder: **Strong evidence** (confidence ≥ 7 and 2+ sources) / **Early signal** / collapsed All findings. Nothing changes; the legend simply explains it to the reader in one line.

Ladder 4 • Priority – *what should we do first?*

Priority is model judgment, and the data shows absolute judgment inflates (4 of 4 "high"). The reliable thing an AI can do is **rank** – so the prompt is changed to demand a strict ranking, and code assigns the words by position:

- Act now** – rank #1 only. Exactly one per update. This is the dashboard's "The one thing" hero card from the spec – same mechanism, now enforced.
- Plan next** – ranks #2-3.
- Worth considering** – the rest.

Scarcity is what makes the word mean something: "Act now" can never appear four times.

What the AI is no longer allowed to write

One rule added to the synthesis prompts (new prompt versions, no pipeline changes): **the narrative describes what and why – never how much.** Banned from prose: intensity and frequency claims – *very, extremely, significantly, overwhelmingly, huge, massive, most, many, widespread, frequently, consistently, growing...* The badges next to the text already say how much, backed by counts.

Still allowed (this is where your "too limited to represent the situation" concern is answered): everything else. Specific consumer language, mechanisms, contrasts the data actually shows ("insurance frustration clusters around Ossur, while price objections cluster around Ottobock"), and grounded recommendations. The prose stays rich – it just stops free-styling the measurements.

Before / after – same finding, "Questions about Ossur products"

Today (uncalibrated): "There is significant interest in Ossur products, with many commenters asking very specific questions." – *significant? many? very? – unfalsifiable, and every reader scales it differently.*

Dominant · 21 of 36 Ossur conversations

54 voices

Curious

Calibrated: "People ask where to buy, what products cost, and whether specific knees fit their amputation level – buying-readiness questions, not idle curiosity. Most go unanswered in the thread." – *the badges carry the magnitude; the prose carries the meaning.*

The legend

A small "How to read this" popover, one component on every data page – the fix for the second half of the problem (different readers, different interpretations). One line per ladder, e.g. "*Dominant = this theme appears in at least 40% of the group's analysed conversations.*" Every reader maps the same word to the same rule.

What this does not change

- The locked spec – this implements its "scores are never displayed as numbers" rule more strictly, and its counts ("number of voices") were already allowed.
- The pipeline – no new passes, no new AI cost. Every ladder is computed in code from fields already persisted per run. Prompt text changes only (B/C/D get the prose rule + the ranking rule).
- Existing data – Monday's run picks up new prompt versions; current pages can render badges from existing data immediately.

Your sign-off – four decisions

1 • The words. Prevalence: **Dominant / Widespread / Recurring / Early signal.** Priority: **Act now / Plan next / Worth considering.** Sentiment: **Strongly positive ... Polarized ... Strongly negative.** These are client-facing vocabulary – veto or reword anything.

2 • Numbers on the badge. Recommended: yes – "Dominant · 21 of 36 Ossur conversations · 54 voices". The alternative is word-only badges with counts in the legend/hover. My call would be numbers-on: it's the anti-sampling proof point.

3 • Forced priority ranking. Exactly one "Act now" per update, assigned by rank, not by the model's own rating. Yes/no.

4 • Thresholds. 40% / 15% / floors as tabled – set from the one real run we have; they're config, reviewed on Monday's output alongside the spec's existing gate-threshold review. Approve as starting values.