

Product Page Redesign — Spec, Draft 1

3 July 2026 · for Heinrich's review · nothing in here is built yet

The problem being fixed: the pages render the database, not the findings. Each data page is a flat dump of one pipeline table, every AI observation gets equal weight, and confidence scores are shown raw — so the client does the extraction work the product is supposed to do. This spec reshapes every page around a client question, with a hard editorial layer on top of the same data. The visual language from the green refresh is unchanged; this is an information & content redesign.

1 · Global rules (apply to every page)

Report = delta, Dashboard = state, pages = depth. The weekly email says *what changed*. The dashboard says *where you stand*. Each data page answers one client question in depth, with evidence one click away. No surface repeats another.

Curation gate. Two tiers, applied everywhere:

Confirmed — confidence ≥ 7 and ≥ 2 evidence sources → shown in the headline zone.

Early signal — below the gate but strength ≥ 6 , or a single-source theme → shown in a clearly-labelled "Early signals" section ("worth watching, not yet confirmed").

Everything else lives in a collapsed "All findings" view. Thresholds tuned on the live Ossur data during build; they are config, not hard-code.

Scores gate and order — they are never displayed as numbers. "conf 4 · opp 6" disappears.

Replacements: a "Strong evidence" chip (gate passed), "Early signal" chip, and sort order. The product expresses judgment; it doesn't hand the client a probability.

Client-facing language ban list. Never shown: *run, pass, gather, scraped, pipeline, corpus, analysed videos, run id*. Replacements: "conversations analysed", "sources", "update", "data through <date>". Every empty state rewritten in client language (e.g. "Your first analysis lands Monday morning" — not "Pass A has not run"). Coverage stats are reframed as a trust signal, never as telemetry.

"New" badge. Any theme or insight that wasn't present in the previous update gets a small "New" marker — the site-side cue for someone who skimmed the email. Powered by the lightweight theme-matching in §8; markers only, never a prose summary (that's the report's job).

2 · Dashboard — "Where do we stand?"

The welcome page: a **state snapshot**, deliberately not this week's news.

1. **Welcome band** — "Ossur — what your market is saying." Coverage line in human terms: "*Tracking 7 search terms across TikTok, YouTube & Instagram · 2,651 conversations analysed · data through 30 Jun · next update Monday.*" Sources: `tracking_configs`, latest run timestamps.
2. **Where you stand** — three stat cards, state-framed: sentiment position (positive share of analysed conversations), share of tracked conversation (client vs each competitor, from `run_summary`), audience mood (dominant emotions from `audience_insights.emotion` — currently collected, surfaced nowhere).
3. **What your market is talking about** — top 3 current themes (by evidence count × strength), one line each on why it matters, each routing to Voice (filtered) or Market. "New" badges apply.
4. **The one thing** — the single top-priority, best-grounded recommendation as the hero card → routes to Market Intelligence. One, not a list.

Removed from this page: top-videos table and content-type tiles (→ Content page), run ids, scraped/analysed counts as KPIs, every pipeline-jargon empty state.

3 · Market Intelligence — "*What should we do?*"

1. **Consumer Intelligence Summary** (new Pass D output, §8) — the headline block: top 3 unmet needs · top 3 buying triggers · top 3 differentiators vs competitors · audience emotional snapshot · threats. This is the "someone already read everything for you" element the page is missing.
 2. **Top recommendations** — promoted above insights (they're the outcome). Max 3, gate-passed, each with its reasoning and a "See the voices behind this" evidence link. Rest collapsed under "All recommendations".
 3. **Key insights** — max ~3, gate-passed, ordered by opportunity. Theme chips + deep-link to Voice stay as built.
 4. **Early signals** — below-gate insights + badged single-source themes, honestly framed.
 5. **All findings** — collapsed archive (essentially today's page, demoted).
-

4 · Voice of Customer — "*What are they actually saying?*"

1. **Themes as the organizing unit** — not the raw insight list. Each theme: clean canonical label (Pass B, §8), synthesized description, emotion, number of voices, expandable verbatim quotes from `insight_evidence`. The existing filter bar stays; add journey-stage filter once the new tag lands.
2. **"How your customers talk"** — the language-samples panel: verbatim phrase clusters for marketing copy (new Pass A output, §8). For a product named Verbatim, this is the signature panel.
3. **Early signals** — single-source themes, badged.

5 • Competitive — "Where do we stand vs Ottobock?"

- 1. **The ownership story** — leads the page: which themes each brand owns, content gaps, threats (curated Pass C output). Qualitative first, per the SOV decision.
- 2. **"What their customers say about you"** — hero slot for cross-reference switching signals (client mentions found under competitor/industry videos, §8). The single most valuable comment type in the corpus, currently lost in Ottobock's bucket.
- 3. **Share of tracked conversation** — the existing bar, kept as supporting evidence below the qualitative story, honest label unchanged.
- 4. **Brand presence read** — the `keyword_performance` diagnostic made client-facing: "the category talks, the brand doesn't" as a narrative card (brand vs competitor term survival + volume). Already computed every run, surfaced nowhere.

Phase 2 (flagged, your call): AI Share of Voice — how ChatGPT / Gemini / Perplexity describe Ossur vs Ottobock. Cheap via the existing LLM stack; slots in as one section here. Not in the first build pass.

6 • Content — "What content works in this niche?"

- 1. **Hook intelligence** — `hook_style` aggregated by engagement, with real `hook_text` examples. Collected for every analysed video since v4.1; currently a hidden table column.
- 2. **Content-type performance** — the tiles, moved here from the dashboard.
- 3. **Top voices** — creator intelligence: the accounts driving conversation in the category (aggregate `account_name`: videos, views, entity mix). Strategy content-pillar #5, never built — it's a group-by away.
- 4. **All videos** — the existing table, as the evidence layer at the bottom.

7 • Reports page & the weekly email

- **Email leads with the delta block** — "what changed since last week": sentiment shift, share shift, new themes count, the top new recommendation. Computed from consecutive `run_summary` rows — no genealogy needed.
- **Every item in the email deep-links** to its page section on the site. The report becomes the login driver; the site never repeats the report's prose.
- Reports page itself: unchanged archive of past reports.

8 • Pipeline changes (the approved v5 pull-forwards)

Change	What it does	Surfaces on
Pass A prompt v3	Adds <code>switching_signal</code> + <code>buying_trigger</code> insight categories, a <code>journey_stage</code> tag, and a <code>language_samples[]</code> output. Prompt-version bump; schema touch is additive.	Market (CI Summary), Voice (language panel, filters)

Cross-reference detection	Regex pass over comments on competitor/industry videos for client-brand mentions → flagged as switching signals attributed to the client.	Competitive ("what their customers say about you")
Step A2 fixes	(a) Single-source themes badged instead of dropped at the evidence floor; (b) cluster at bucket level so one concern split across categories re-stitches into one theme.	Early signals everywhere; stronger themes on Voice
Pass B	Canonical theme labels + descriptions (~\$0.02/run). Labels are now page headlines, so polish is front-of-house.	Voice, Dashboard themes
Pass D split + CI Summary	D-a insights / D-b recommendations grounded via evidence retrieval (the fix for generic recs — the trigger fired on the Ossur run), plus the <code>consumer_intelligence_summary</code> output.	Market Intelligence lead + recommendations
Mini theme-matching	Embedding-match latest-run themes against the previous run → <code>first_seen</code> flag. Powers "New" badges + the email delta. Deliberately <i>not</i> full genealogy.	All pages, weekly email

Note: existing Ossur/Sealand data won't carry the new fields until the next scheduled run — every new section must degrade gracefully (hide, never show empty jargon). Honest caveat: curation and grounding can't conjure depth from a thin corpus; on Ossur, pages will be sparse-but-true. That's the correct trade at this price point — and the real fix for depth is a data-source decision (below), not presentation.

9 · IA & housekeeping

- **Sidebar:** Dashboard · Market Intelligence · Voice of Customer · Competitive · Content · Reports, then a separated bottom group: Team · Billing · Settings. Trends stays hidden.
- **Settings split (facts vs knobs):** client-editable = company name, competitor names, report emails/day. Superadmin-only (hidden from clients) = keywords, max videos/comments, comment depth, platforms. Comment depth is a pricing lever, not a user preference.

10 · Flagged decisions — yours, not built until called

Reddit (V2). The highest-value item if effort is ignored — it attacks comment depth, the constraint capping insight quality everywhere. Pages in this spec are written platform-agnostic so Reddit slots in without rework. Separate go/no-go.

AI Share of Voice — Competitive page phase 2 (\$5).

Onboarding auto-suggest (company name → proposed keywords/competitors, client confirms) — the facts-not-expertise fix for self-serve. Separate small follow-on after the pages.

Gate thresholds — confirmed/early-signal cutoffs reviewed together on real Ossur output after the first build pass.

11 · Build order

1. **Pre-req (operational):** merge `feat/per-keyword-search` before Monday's scheduled run, and check the URL-overflow pagination fix on the Inngest path.
2. Pipeline changes (§8) — so Monday's run already produces the new outputs.
3. Market Intelligence → Dashboard → Voice → Competitive → Content (one page per pass, your review between pages, co-design style).
4. Weekly email: delta block + deep links.
5. Settings split + language-ban sweep across all empty states.

Draft 1 · 2026-07-03 · react to anything — sections, ordering, gates, naming. Nothing here is committed until you say so.